

The background features a large, abstract graphic on the left side composed of numerous overlapping blue triangles of varying sizes and shades, creating a sense of depth and perspective. The right side of the slide is a solid light gray.

ZOOM BEHAVIOR INSIGHT

Presented by:
Ofir Duek , Aviv Meir
, Rotem Aloni

PROJECT REVIEW

Problem

In online classes, instructors can't reliably infer students' behaviors from a single webcam feed (e.g., gaze, privacy, headphones, or object usage).

Our Goal

Given a single Zoom-like webcam frame, predict a discrete behavior vector:
Gaze, Headphones, Environment, Privacy, Object-in-hand (incl. Unknown).

What's new since the first presentation:

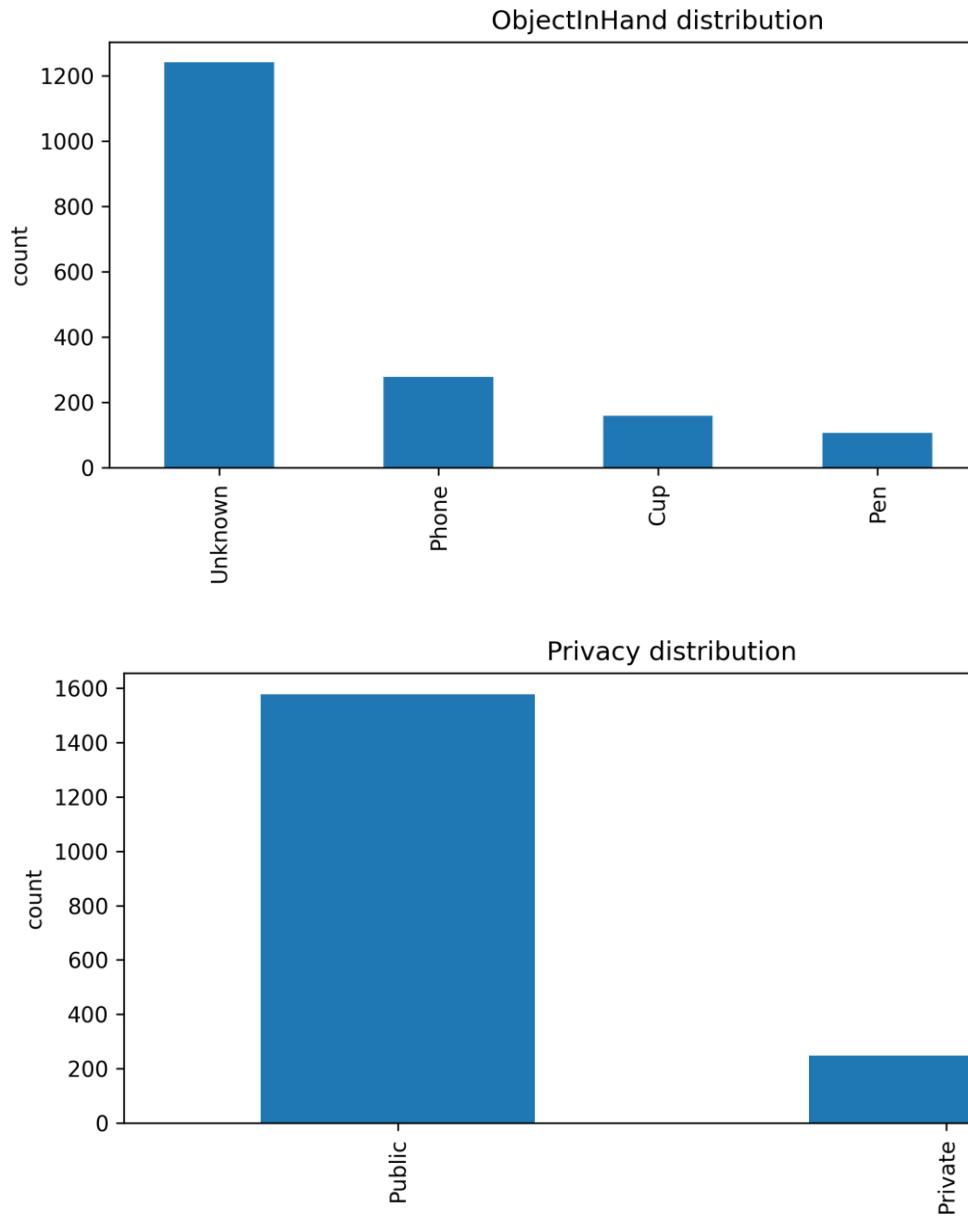
- Switched from 0 - 100% statistics to categorical behavior labels.
- Set strict labeling rules with Unknown when evidence is missing.
- Started the data pipeline: real Zoom-like samples + synthetic augmentation.

Novelty / Contribution

We output multiple interpretable behaviors (not a single engaged score) and use controlled synthetic augmentation to cover rare cases.

RELATED WORK

| Paper (Year) | Task | Signals / Cues | Method / Model | Data | How it relates to us |
|---|---------------------------|--------------------------------|--|-------------------------------|---|
| Stungage (2022) ^[1] | Engagement monitoring | Webcam video + meeting context | Real-time system + ML pipeline | Online meeting sessions | We predict a multi-attribute behavior vector from a single Zoom-like frame (not see engagement score). |
| Facial behavior engagement (2022/2023) ^[2] | Engagement estimation | Face cues (gaze, head pose) | Sequence model (LSTM on face features) | Online learning webcam videos | Similar gaze setting, but we output multiple behavior labels (incl. Unknown) and add non-face cues. |
| Emotion + eye/head (2019) ^[3] | Engagement classification | Emotion + eye/head movement | ML classifier on multi-cue features | Webcam e-learning sessions | Multi-cue idea fits; we extend cues to privacy / environment / object-in-hand and handle missing evidence with Unknown. |



DATASET OVERVIEW + EDA

- Our dataset combines real Zoom-like photos and synthetic images.
 - Each image is labeled with a 5-attribute discrete behavior vector, including Unknown when evidence is missing.
 - EDA: class distributions per attribute to identify imbalance and guide augmentation.
-

REAL DATA

- Real Zoom-like images from team-captured photos and public online datasets.
- Used as the most realistic reference for training and evaluation.

Labels:

Gaze = Not_Camera,
Headphones =
With_Headphones,
Environment = Outdoor,
Privacy = Private ,
ObjectInHand = Unknown



Labels:

Gaze = Not_Camera,
Headphones =
Without_Headphones,
Environment = Outdoor,
Privacy = Public,
ObjectInHand = pen

STABLE DIFFUSION (TEXT-TO-IMAGE)

- Generated images from scratch using prompts.
 - Used to control behaviors and increase rare cases.
 - Sample prompts:
 1. Zoom-style webcam photo, gaze at camera, with headphones, outdoor public background with blurred people, holding a pen in hand. Realistic, natural daylight, shallow depth of field, sharp subject.
 2. Zoom-style webcam photo, eyes closed, with headphones, indoor private setting, holding a cup in hand. Realistic, soft light, shallow depth of field, centered framing.
-



INPAINTING BACKGROUND REPLACEMENT

- Replace the background only to control Environment and Privacy labels.
- Keep the same subject to reduce confounding factors.
- Produces paired samples: Before → After with consistent appearance.

Same subject, different context

Before



After



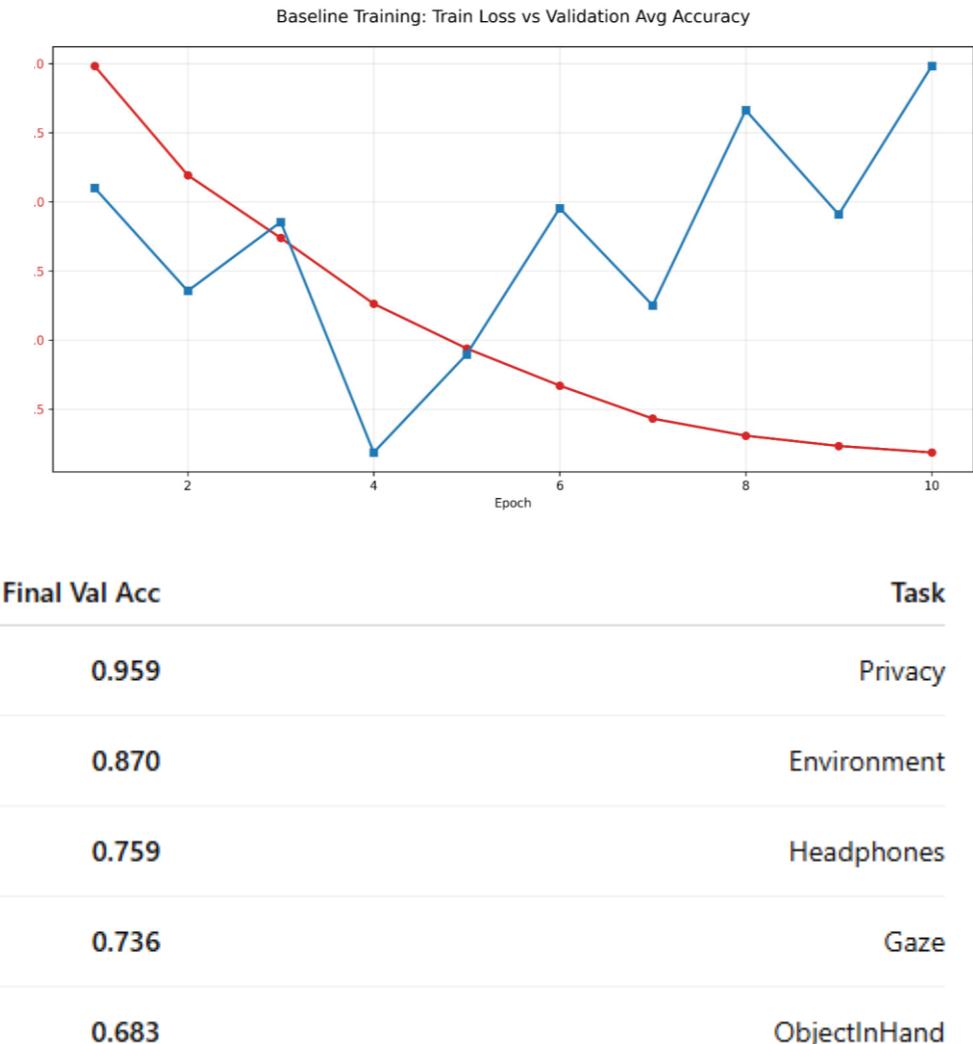
BASELINE MODEL (RESNET18 MULTI-TASK) - RESULTS

Baseline setup:

- **Model:** ResNet18 (ImageNet pretrained), with one classification head per behavior.
- **Training:** 10 epochs, batch=32, Adam lr=1e-3, 80/20 train–val split.

Key takeaway:

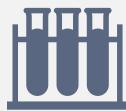
- Train loss decreases strongly while validation avg accuracy changes only mildly → baseline overfitting.
- Validation performance is uneven across tasks (best: Privacy, hardest: ObjectInHand).



NEXT STEPS



Train an improved model on the finalized dataset (stronger baseline + regularization).



EDA-driven data balancing: add targeted real/synthetic samples for rare classes.



Evaluate on a held-out real test set and report per-behavior metrics.



Error analysis + prototype demo: input one Zoom-like frame → output the behavior vector.